

LLM Model Pricing Comparison

July 2026 - prices per million tokens (MTok), USD

MODEL	PROVIDER	INPUT	CACHED INPUT	OUTPUT
Claude Fable 5	Anthropic	\$10.00	\$1.00	\$50.00
Claude Opus 4.8	Anthropic	\$5.00	\$0.50	\$25.00
Claude Sonnet 5	Anthropic	\$3.00	\$0.30	\$15.00
Claude Sonnet 5 Intro pricing until Aug 31, 2026	Anthropic	\$2.00	\$0.20	\$10.00
GPT-5.5	OpenAI	\$5.50	\$0.55	\$33.00
GLM 5.2	Zhipu AI	\$1.54	\$0.154	\$4.84

Key Takeaways

Input volume usually drives the bill, but output is where you save. Output is priced 5-6x higher per token and is directly controllable. Shorter responses, structured extraction, and tool use over prose can cut cost significantly.	GLM 5.2 is dramatically cheaper. At \$4.84/MTok output, GLM 5.2 costs roughly one-third of Sonnet 5 and one-fifth of GPT-5.5, while claiming Opus-level coding capability.
Fable 5 is the frontier premium. At \$50/MTok output, Fable 5 is 2x Opus 4.8 and 10x Sonnet 5. Use it only when the quality gap justifies the cost.	Sonnet 5 introductory pricing expires August 31, 2026. Until then, Sonnet 5 at \$2/\$10 is roughly cost-neutral with Sonnet 4.6 despite the new tokenizer producing about 30% more tokens for the same text.

Why Anthropic's New Tokenizer Matters

Starting with Opus 4.7, Anthropic introduced a new tokenizer used by all subsequent models: Opus 4.8, Fable 5, and Sonnet 5. It produces **about 30% more tokens for the same text** compared to earlier models like Sonnet 4.6 and Opus 4.6.

How we know: Opus 4.6 fits about 750K words per 1M tokens. Opus 4.7+ fits about 555K words per 1M tokens. That is 35% fewer words per token, or equivalently, about 30% more tokens for the same input.

What it means in practice:

- A workload that cost \$100/month on Sonnet 4.6 becomes about \$130/month on Sonnet 5 at the same per-token rate, even though the per-token price did not change.
- Anthropic's introductory pricing (\$2/\$10 vs. standard \$3/\$15) is designed to offset this: the lower per-token price roughly cancels out the higher token count.
- The old Sonnet 4.6 and Opus 4.6 still use the previous tokenizer, so they remain cheaper per word for identical content.

The same problem exists across providers. OpenAI and Zhipu use their own tokenizers, which also differ from each other and from Anthropic. Per-token prices do **not** reflect actual cost for identical workloads, so benchmark on your own data.

When comparing Anthropic models to each other, compare by cost-per-word, not cost-per-token. **When comparing across providers**, the tokenizer gap makes price-per-token even less reliable.

How to Choose

Choose the model based on task complexity, not per-token price. A cheap model that fails on a hard task wastes more tokens and time than an expensive model that succeeds in one pass.

Use case	Recommended model
Routine tasks, bulk extraction, simple Q&A	Sonnet 5 or GLM 5.2
Hard reasoning, agentic coding, complex workflows	Opus 4.8
Frontier quality at any cost	Fable 5

The cheapest token is the one you never send. Before calling the model, ask: can a script do this? Can a template replace this prompt? Can cached results be reused? Can the output be shorter? Every token saved is a token you do not pay for, at any price tier.

Sources: Anthropic pricing page (July 2026), OpenAI API pricing (July 2026), Fireworks AI serverless pricing (July 2026). All prices shown include 10% Data Zone uplift where applicable.